

Machine Learning Shiny App

Vin Seixas, Samuel de Oliveira, Hunter Stopford

Github link: <https://github.com/viniciusdel/MachineLearningPredictions>

Abstract

In this project, we apply various Machine Learning (ML) models to quintessential problems in the field of computational statistics, and create a user-interactive Shiny app for demonstration purposes. A Random Forest (RF) classifier was used on a large dataset of handwritten numerical digits in order to train a model to recognize numbers rendered on a digital draw-pad. We then applied a Support Vector Machine (SVM) model to a medical cost dataset to create a regression model which could then be used in order to make predictions of insurance costs for the user, given relevant data like age, sex, BMI, region, etc. Lastly, we create a medical questionnaire which uses a RF classification model from the caret package to make a highly-predictive diabetes diagnosis for users. We discuss the differences between the datasets and the metrics used to evaluate the ML algorithms. Lastly, we explain the design of our Shiny application and how the user interface was integrated into the trained ML models.

Introduction

We tested our datasets (Digit Recognition notwithstanding) across a variety of ML algorithms, including Classification and Regression Trees (CART), Linear Discriminant Analysis (LDA), K-Nearest Neighbor (KNN), Neural Networks (NN), and Random Forest. The selection of model we would use in our shiny application was based on which produced the least error after training. The particular R packages and ML algorithms we used for each tab are as follows:

- Digit Recognition - randomForest package - Random Forest
- Insurance Cost - caret package - Support Vector Machine
- Diabetes Diagnosis - caret package - Random Forest

We chose for our datasets one regression example--insurance cost prediction, and two classification examples--digit recognition and diabetic diagnosis. Broadly speaking, supervised machine learning (more on supervised vs. unsupervised in the Discussion section) can be divided into 'regression' and 'classification' algorithms, where a distinguishing feature is whether the output variable type is of a continuous or discrete value. Regression algorithms try to make the best fit line between input variables and values which can be real and continuous, like the price of insurance. Classification algorithms try to make a decision boundary such that output value fits into a discrete value or class, like the binary logic of a diagnosis or the value of an integer.

Methods

Digit Recognition

For the digit recognition application, we trained an RF model with the MNIST database which is often used for testing and validating optical character recognition devices. The model is trained on a dataset of 42,000 known observations and then validated against another 28,000 observations. See table below:


Number of Observations for Each Digit

0	1	2	3	4	5	6	7	8	9
4132	4684	4177	4351	4072	3795	4137	4401	4063	4188

The actual arguments fed into the algorithm are grayscale pixel values from 0 to 255, which form a 28x28 image, resulting in a total of 784 pixels.

One important parameter selected was the 'number of trees' which determines the number of decision trees the algorithm should build prior to evaluating the predictions of each tree, and determining which class was selected by the most number of trees. For our model, we chose 25 as the number of trees. A greater number of trees will also make the model training time slower. The training time given our parameters was around 209.9 seconds.

The number of trees was decided based on flattening trend-line in the error reduction as the number approached 25. We can see the error decreases as we add more trees, but anything more than 25 would not significantly increase accuracy:



The data frame also returns a confusion matrix generated after testing. The confusion matrix shows the number of times a digit is correctly identified, and how often it is confused for each of the other classes. We can see the results of the trained model in the following confusion matrix:

```

Type of random forest: classification
Number of trees: 25
No. of variables tried at each split: 28

OOB estimate of error rate: 6.3%
Confusion matrix:

```

	0	1	2	3	4	5	6	7	8	9	class.error
0	4027	1	11	8	4	15	27	4	24	11	0.02541142
1	1	4588	18	21	9	10	6	6	14	11	0.02049530
2	20	13	3913	42	33	17	30	45	49	15	0.06320326
3	8	9	101	3940	7	118	13	44	76	35	0.09446104
4	11	11	20	9	3816	13	28	6	21	137	0.06286837
5	20	12	17	124	15	3479	40	12	45	31	0.08326746
6	33	10	19	3	27	43	3982	1	18	1	0.03746676
7	7	17	68	20	34	7	1	4150	16	80	0.05681818
8	18	34	53	79	36	75	31	15	3658	64	0.09968004
9	29	9	19	48	129	28	8	66	50	3802	0.09216810

We were able to train the model to predict numbers with only a 6.3% error.

Insurance Cost

In contrast to the previous algorithm, predicting the cost of insurance premiums would require a regression algorithm to generate a continuous value output. Our dataset was composed of 1,338 observations with 7 factors each, 80% of which was used for training and 20% for testing. A regression fit was attempted using CART, LDA, KNN, SVM, RF, and NN. LDA and NN failed to converge entirely, which may be due to formatting issues in the data-set. The metric used to evaluate the suitability of any particular matrix was the Mean Absolute Error and the Root Mean Squared Error. Although, RF and SVM had similar results, with some tweaking SVM produced slightly better results. See table and plot below:

```
Call:
summary.resamples(object = results)
```

```
Models: cart, knn, svm, rf, nnet
Number of resamples: 10
```

MAE


	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
cart	2681.995	3218.822	3472.598	3706.852	4143.310	4980.824	0
knn	7146.108	7485.428	7799.340	7882.858	8356.537	8605.010	0
svm	1836.417	2446.088	2556.530	2564.770	2782.064	3047.852	0
rf	1972.843	2396.621	2641.320	2573.146	2758.312	3066.249	0
nnet	12632.847	12934.084	13215.645	13187.160	13453.747	13703.147	0

RMSE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
cart	3721.808	4720.339	5237.334	5308.145	5847.030	6950.018	0
knn	9724.608	10543.973	10808.022	10906.570	11079.246	12483.971	0
svm	3104.436	4315.990	4641.125	4652.820	5127.560	5804.294	0
rf	2937.678	4184.731	4708.914	4550.633	4983.398	5648.545	0
nnet	16900.679	17113.729	17718.930	17767.779	18147.124	19276.253	0

Rsquared

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
cart	0.61445255	0.7737296	0.8234133	0.7968216	0.8357828	0.9138882	0
knn	0.09233985	0.1441629	0.1630478	0.1746007	0.2089587	0.2731238	0
svm	0.78568037	0.8111291	0.8536403	0.8486732	0.8725954	0.9410740	0
rf	0.78948683	0.8257163	0.8504202	0.8531205	0.8737969	0.9542256	0
nnet	NA	NA	NA	NaN	NA	NA	10



Diabetes

In the diabetes prediction application, we went a different route and used the *caret* package and its Random Forest classifier to train a model on patient data of people that suffer from diabetes. 520 patients were analyzed, with fields ranging from Age and Gender to symptoms like irritability. 200 of those patients were negative for diabetes, while the rest were positive.

We analyzed different models like CART, LDA, K nearest neighbors, random forests and neural networks, but you can see that random forests had the highest accuracy mean, so that's what we picked.

	Overall
PoluryaYes	100.000000
PolydipsiaYes	83.294493
Age	27.624015
GenderMale	26.684665
IrritabilityYes	9.408118
AlopeciaYes	8.487711
DelayedHealingYes	8.416027
suddenWeightLossYes	8.095202
PartialParesisYes	7.082277
ItchingYes	4.231519

Accuracy								
	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's	
cart	0.7619048	0.8302846	0.8554007	0.8461672	0.8750000	0.9024390	0	
lda	0.8048780	0.8571429	0.8571429	0.8724739	0.8963415	0.9523810	0	
knn	0.7142857	0.7804878	0.8214286	0.8220093	0.8750000	0.9268293	0	
rf	0.9268293	0.9523810	0.9759001	0.9734611	1.0000000	1.0000000	0	
nnet	0.9024390	0.9515099	0.9642857	0.9638211	0.9761905	1.0000000	0	
Kappa								
	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's	
cart	0.5183486	0.6592883	0.6889002	0.6780094	0.7250288	0.7902813	0	
lda	0.6076555	0.7133300	0.7174888	0.7432124	0.7841767	0.8990385	0	
knn	0.4349776	0.5724218	0.6435685	0.6456811	0.7423176	0.8512696	0	
rf	0.8479604	0.8996310	0.9485360	0.9442979	1.0000000	1.0000000	0	
nnet	0.7950000	0.8992177	0.9249371	0.9240457	0.9501188	1.0000000	0	

An interesting analysis we also did was to see which fields had the highest relevance to the algorithm's prediction. We performed a variable importance analysis that shows that some features are more important than others namely Polyuria and Polydipsia.

Shiny Application

On the application side, the packages used were Shiny, Shinythemes for changing the UI appearance, Magick for image processing and manipulation, Gtools, and Imager for additional image processing functionalities. The Shiny package divides the structure of the application into a 'user interface' and a 'server' side. On the UI side, we organized each of our trained models into separate tabs, along with the user functions for those models. The syntax for the UI design is something like a mix between HTML and R. Elements of a page can be generated with a function call like in R, and each function will have its own function parameters, but specific elements can be placed within one another by separation with a comma. Since the UI has to interact with the server side, variables which must be passed to the server side are identified with 'inputID.' On the server side, those variables are then just referenced using the syntax 'input\$[insert input ID].' We include our already trained models

in a subdirectory of the main file. Those can be referenced and used using the 'readRDS' function.

Conclusion & Future Work

One commonality shared by all the data we used was that it was labeled data, suitable for supervised machine learning. With supervised machine learning, it's necessary to train a model with observations wherein all relevant variables are known. The model is then validated with test data to determine how well it makes a regression fit or classification. Unsupervised machine learning uses unlabeled data, and the algorithm is then left to make associations amongst the variables. This has the added benefit of finding hidden correlations, but also may not offer good results and is more difficult to validate. For future work, we would like to try some unsupervised algorithms on larger, less understood datasets.

One of the challenges we faced in this project was generating and processing the plot in such a way as to make it usable for the digit recognition model. The model seems to be highly sensitive to the order in which we carried out the image processing functions. Certain numbers also reliably produced errors beyond what we saw from the training data, even though the pixel image being fed into the model greatly resembled what we saw in the training data. For future work, we would invest more time into the image processing to reduce the amount of error the draw-pad function produced.

References

M M Faniful Islam, Rahatara Ferdousi, Sadikur Rahman, Yasmin Bushra. UC Irvine. Early Stage Diabetes Risk Prediction dataset.

<https://archive.ics.uci.edu/ml/datasets/Early+stage+diabetes+risk+prediction+dataset>.

Yann LeCun, Corinna Cortes, Christopher J.C. Burges. NIST. MNIST database of handwritten digits.<http://yann.lecun.com/exdb/mnist/index.html>

Marc Agenis. Plotting function on Number Recognition app.

<https://stackoverflow.com/posts/48442522/timeline>